


CS498 AML HW10

Huiyun Wu (hwu63)

Yidi Yang (yyang160)

Page 1: (MNIST with TUTORIAL)

1. Accuracy plot (train and val, with smoothing = 0.41)



Where the orange line represents **train accuracies** (logged every batch), and the red smooth line represents the **validation set accuracy** (logged every 100 batches).

2. Accuracies (train, test, val) from the 'best' model with a brief explanation.

Training set: ~0.995

Testing set: 0.9877

Validation set: 0.9859


The model we use here is the original one from the MNIST tutorial. We trained it with 2001 batches and reached this accuracy.

Page 2: (MNIST w MODIFIED)

1. Description of changes

We added another convolutional layer (with filter = 32 originally to 4) and removed all the poolings, so that the model has three layers in total with filter equals 4 and no max poolings.

2. Accuracy plot (train and val)



3. Accuracies (train, test, val) from the 'best' model with a brief explanation.

Training set: 0.9971

Testing set: 0.985


Validation set: 0.982

The model here is modified as explained above. We runned it for 2001 batches. The accuracy here dropped by 1% compared with the original model, we think this is due to the filter change.


Page 3: (CIFAR w TUTORIAL)

1. Accuracy plot (train and val)

Train:



Val:



2. Accuracies (train, test, val) from the 'best' model with a brief explanation.

We runned the model for around 3900 batches. We don't have GPU so that we couldn't run the model for sufficiently many batches. But we believe the accuracy would be around the desired level (0.83+) if running it with more batches.

Training set accuracy: 0.7117

Testing set (Evaluation set) accuracy: 0.647

Validation set accuracy: 0.6924


Page 4: (CIFAR w MODIFIED)

1. Description of changes


We added another convocational layer between conv2 and normalization 2.

2. Accuracy plot (train and val)

Train:



Val:



3. Accuracies (train, test, val) from the 'best' model with a brief explanation.

Training set accuracy: 0.7330

Testing set accuracy: 0.6769

Validation set accuracy: 0.7070

To make a valid and fair comparison, we also runned it for 3900 batches. The accuracies of three datasets all increased. Because we added another convolutional layer in the model.

1. Creating the model/graph: Corresponds to the `cnn_model_fn` in the mnist tutorial.

```

52         tf.summary.scalar('stddev', stddev)
53         tf.summary.scalar('max', tf.reduce_max(var))
54         tf.summary.scalar('min', tf.reduce_min(var))
55         tf.summary.histogram('histogram', var)
56
57     conv1 = tf.layers.conv2d(
58         inputs=image_shaped_input,
59         filters=4,
60         kernel_size=[5, 5],
61         padding="same",
62         activation=tf.nn.relu)
63
64     conv2 = tf.layers.conv2d(
65         inputs=conv1,
66         filters=4,
67         kernel_size=[5, 5],
68         padding="same",
69         activation=tf.nn.relu)
70
71     conv3 = tf.layers.conv2d(
72         inputs=conv2,
73         filters=4,
74         kernel_size=[5, 5],
75         padding="same",
76         activation=tf.nn.relu)
77
78     conv3_flat = tf.reshape(conv3, [-1, 7 * 7 * 64])
79
80     dense = tf.layers.dense(inputs=conv3_flat, units=1024, activation=tf.nn.relu)
81
82
83     with tf.name_scope('dropout'):
84         keep_prob = tf.placeholder(tf.float32)
85         tf.summary.scalar('dropout_keep_probability', keep_prob)
86         dropout = tf.layers.dropout(inputs=dense, rate=keep_prob)

```

2. Defining loss & optimizer: Part where you set the loss (Corresponds to line 100 in the mnist tutorial) and optimizer.

```

91     with tf.name_scope('cross_entropy'):
92         # The raw formulation of cross-entropy,
93         #
94         # tf.reduce_mean(-tf.reduce_sum(y_ * tf.log(tf.softmax(y)),
95         #                               reduction_indices=[1]))
96         #
97         # can be numerically unstable.
98         #
99         # So here we use tf.losses.sparse_softmax_cross_entropy on the
100        # raw logit outputs of the nn_layer above, and then average across
101        # the batch.
102        with tf.name_scope('total'):
103            cross_entropy = tf.losses.sparse_softmax_cross_entropy(
104                labels=y_, logits=y)
105            tf.summary.scalar('cross_entropy', cross_entropy)

```

3. Training: Part where you either manually defined a train loop or where you used a built-in train function to train the model.

```

159    for i in range(2001):
160
161        if i % 100 == 0: # Record summaries and test-set accuracy
162            print("running test and val...")
163            summary, acc = sess.run([merged, accuracy], feed_dict=feed_dict(0))
164            test_writer.add_summary(summary, i)
165            print('Test Set Accuracy at step %s: %s' % (i, acc))
166            summary2, acc2 = sess.run([merged, accuracy], feed_dict=feed_dict(2))
167            val_writer.add_summary(summary2, i)
168            print('Val Set Accuracy at step %s: %s' % (i, acc2))
169
170        else: # Record train set summaries, and train
171            summary, _ = sess.run([merged, train_step], feed_dict=feed_dict(1))
172            train_writer.add_summary(summary, i)
173            train_writer.close()
174            test_writer.close()
175            val_writer.close()

```

4. Any other snippet you found relevant

```
131 train_images, val_images, train_labels, val_labels = \
132     train_test_split(mnist.train.images, mnist.train.labels, test_size=0.2)
133
134 def next_batch(num, data, labels):
135     """
136     Return a total of `num` random samples and labels.
137     """
138     idx = np.arange(0, len(data))
139     np.random.shuffle(idx)
140     idx = idx[:num]
141     data_shuffle = [data[i] for i in idx]
142     labels_shuffle = [labels[i] for i in idx]
143
144     return np.asarray(data_shuffle), np.asarray(labels_shuffle)
145
146 def feed_dict(train):
147     """Make a TensorFlow feed_dict: maps data onto Tensor placeholders."""
148     if train == 1 or FLAGS.fake_data:
149         xs, ys = next_batch(100, train_images, train_labels)
150         k = FLAGS.dropout
151     elif train == 0:
152         xs, ys = mnist.test.images, mnist.test.labels
153         k = 1.0
154     else:
155         xs, ys = val_images, val_labels
156         k = 1.0
157     return {x: xs, y: ys, keep_prob: k}
```

Validation data set splitting and modified feed_dict.

1. Creating the model/graph: Corresponds to the `cnn_model_fn` in the mnist tutorial.

```

187 def inference(images):
188     """Build the CIFAR-10 model..."""
189     ...
202     with tf.variable_scope('conv1') as scope:
203         kernel = _variable_with_weight_decay('weights',
204                                             shape=[5, 5, 3, 64],
205                                             stddev=5e-2,
206                                             wd=None)
207         conv = tf.nn.conv2d(images, kernel, [1, 1, 1, 1], padding='SAME')
208         biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.0))
209         pre_activation = tf.nn.bias_add(conv, biases)
210         conv1 = tf.nn.relu(pre_activation, name=scope.name)
211         _activation_summary(conv1)
212
213     # pool1
214     pool1 = tf.nn.max_pool(conv1, ksize=[1, 3, 3, 1], strides=[1, 2, 2, 1],
215                             padding='SAME', name='pool1')
216
217     # norm1
218     norm1 = tf.nn.lrn(pool1, 4, bias=1.0, alpha=0.001 / 9.0, beta=0.75,
219                       name='norm1')
220
221     # conv2
222     with tf.variable_scope('conv2') as scope:
223         kernel = _variable_with_weight_decay('weights',
224                                             shape=[5, 5, 64, 64],
225                                             stddev=5e-2,
226                                             wd=None)
227         conv = tf.nn.conv2d(norm1, kernel, [1, 1, 1, 1], padding='SAME')
228         biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.1))
229         pre_activation = tf.nn.bias_add(conv, biases)
230         conv2 = tf.nn.relu(pre_activation, name=scope.name)
231         _activation_summary(conv2)
232
233     with tf.variable_scope('conv3') as scope:
234         kernel = _variable_with_weight_decay('weights',
235                                             shape=[5, 5, 64, 64],
236                                             stddev=5e-2,
237                                             wd=None)
238         conv = tf.nn.conv2d(conv2, kernel, [1, 1, 1, 1], padding='SAME')
239         biases = _variable_on_cpu('biases', [64], tf.constant_initializer(0.1))
240         pre_activation = tf.nn.bias_add(conv, biases)
241         conv3 = tf.nn.relu(pre_activation, name=scope.name)
242         _activation_summary(conv3)
243
244     # norm2
245     norm2 = tf.nn.lrn(conv3, 4, bias=1.0, alpha=0.001 / 9.0, beta=0.75,
246                       name='norm2')
247
248     # pool2
249     pool2 = tf.nn.max_pool(norm2, ksize=[1, 3, 3, 1],
250                                 strides=[1, 2, 2, 1], padding='SAME', name='pool2')

```

2. Defining loss & optimizer: Part where you set the loss (Corresponds to line 100 in the mnist tutorial) and optimizer.

```

95     # Calculate loss.
96     loss = cifar10.loss(logits, labels)

```

```

334 def train(total_loss, global_step):
335     """Train CIFAR-10 model..."""
336     # Variables that affect learning rate.
337     num_batches_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_TRAIN / FLAGS.batch_size
338     decay_steps = int(num_batches_per_epoch * NUM_EPOCHS_PER_DECAY)
339
340     # Decay the learning rate exponentially based on the number of steps.
341     lr = tf.train.exponential_decay(INITIAL_LEARNING_RATE,
342                                   global_step,
343                                   decay_steps,
344                                   LEARNING_RATE_DECAY_FACTOR,
345                                   staircase=True)
346     tf.summary.scalar('learning_rate', lr)
347
348     # Generate moving averages of all losses and associated summaries.
349     loss_averages_op = _add_loss_summaries(total_loss)
350
351     # Compute gradients.
352     with tf.control_dependencies([loss_averages_op]):
353         opt = tf.train.GradientDescentOptimizer(lr)
354         grads = opt.compute_gradients(total_loss)
355
356     # Apply gradients.
357     apply_gradient_op = opt.apply_gradients(grads, global_step=global_step)
358
359     # Add histograms for trainable variables.
360     for var in tf.trainable_variables():
361         tf.summary.histogram(var.op.name, var)
362
363     # Add histograms for gradients.
364     for grad, var in grads:
365         if grad is not None:
366             tf.summary.histogram(var.op.name + '/gradients', grad)
367
368     # Track the moving averages of all trainable variables.
369     variable_averages = tf.train.ExponentialMovingAverage(
370         MOVING_AVERAGE_DECAY, global_step)
371     with tf.control_dependencies([apply_gradient_op]):
372         variables_averages_op = variable_averages.apply(tf.trainable_variables())
373
374     return variables_averages_op

```

3. Training: Part where you either manually defined a train loop or where you used a built-in train function to train the model.

```

105 class _LoggerHook(tf.train.SessionRunHook):
106     """Logs loss and runtime."""
107
108     def begin(self):
109
110     def before_run(self, run_context):
111         self._step += 1
112         return tf.train.SessionRunArgs(loss) # Asks for loss value.
113
114     def after_run(self, run_context, run_values):
115         if self._step % FLAGS.log_frequency == 0:
116             current_time = time.time()
117             duration = current_time - self._start_time
118             self._start_time = current_time
119
120             loss_value = run_values.results
121             examples_per_sec = FLAGS.log_frequency * FLAGS.batch_size / duration
122             sec_per_batch = float(duration / FLAGS.log_frequency)
123
124             format_str = ('%s: step %d, loss = %.2f (%.1f examples/sec; %.3f '
125                           'sec/batch)')
126             print(format_str % (datetime.now(), self._step, loss_value,
127                               examples_per_sec, sec_per_batch))
128
129     with tf.train.MonitoredTrainingSession(
130         checkpoint_dir=FLAGS.train_dir,
131         hooks=[tf.train.StopAtStepHook(last_step=FLAGS.max_steps),
132               tf.train.NanTensorHook(loss),
133               _LoggerHook()],
134         config=tf.ConfigProto(
135             log_device_placement=FLAGS.log_device_placement)) as mon_sess:
136         while not mon_sess.should_stop():
137             mon_sess.run(train_op)
138
139

```

4. Any other snippet you found relevant


```

208 def inputs(eval_data, data_dir, batch_size):
209     """Construct input for CIFAR evaluation using the Reader ops.
210
211     Args:
212         eval_data: bool, indicating if one should use the train or eval data set.
213         data_dir: Path to the CIFAR-10 data directory.
214         batch_size: Number of images per batch.
215
216     Returns:
217         images: Images. 4D tensor of [batch_size, IMAGE_SIZE, IMAGE_SIZE, 3] size.
218         labels: Labels. 1D tensor of [batch_size] size.
219     """
220     # if not eval_data:
221     #     filenames = [os.path.join(data_dir, 'data_batch%d.bin' % i)
222     #                 for i in xrange(1, 6)]
223     #     num_examples_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_TRAIN
224     if eval_data == 1:
225         filenames = [os.path.join(data_dir, 'data_batch%d.bin' % i)
226                     for i in range(1, 5)]
227         num_examples_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_TRAIN
228     elif eval_data == 2:
229         filenames = [os.path.join(data_dir, 'data_batch%d.bin' % i)
230                     for i in range(5, 6)]
231         num_examples_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_VAL
232     else:
233         filenames = [os.path.join(data_dir, 'test_batch.bin')]
234         num_examples_per_epoch = NUM_EXAMPLES_PER_EPOCH_FOR_EVAL

```

Modified inputs in cifar_input.py to include the validation.